



Unifying Big Data Workloads in Apache Spark

Hossein Falaki
@mhfalaki



Outline

- What's Apache Spark
- Why Unification
- Evolution of Unification
- Apache Spark + Databricks
- Q & A




What's Apache Spark

What is Apache Spark?

General cluster computing engine
that extends MapReduce

Rich set of APIs and libraries

Large community: 1000+ orgs,
clusters up to 8000 nodes





Why Unification?

Unique Thing about Spark

Unification: same engine and same API for diverse use cases

- Streaming, batch, or interactive
- ETL, SQL, machine learning, or graph
- Scala, Python, R, and SQL

This talk: Why unification? And how has it evolved in Spark?

Why Unification?

MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat

jeff@google.com, sanjay@google.com

Google, Inc.

Abstract

MapReduce is a programming model and an associated implementation for processing and generating large data sets. Users specify a *map* function that processes a key/value pair to generate a set of intermediate key/value pairs, and a *reduce* function that merges all intermediate values associated with the same intermediate key. Many

given day, etc. Most such computations are conceptually straightforward. However, the input data is usually large and the computations have to be distributed across hundreds or thousands of machines in order to finish in a reasonable amount of time. The issues of how to parallelize the computation, distribute the data, and handle failures conspire to obscure the original simple computation with large amounts of complex code to deal with


Why Unification?

MapReduce: a general engine for batch processing

We wrote the first version of the MapReduce library in February of 2003, and made significant enhancements to it in August of 2003, including the locality optimization, dynamic load balancing of task execution across worker machines, etc. Since that time, we have been pleasantly surprised at how broadly applicable the MapReduce library has been for the kinds of problems we work on. It has been used across a wide range of domains within Google, including:

Big Data Systems Today

MapReduce




Pregel Giraph
Dremel Millwheel
Storm Impala
Drill S4 ...

General batch
processing

Specialized systems
for new workloads


Hard to *combine* in pipelines

Big Data Systems Today




Key Idea

MapReduce + data sharing (RDDs) captures most distributed apps



Iterative



Streaming

Benefits of Unification

1. Simpler to **use** and **operate**
2. Code **reuse**: e.g. only write monitoring, FT, etc once
3. New **apps** that span processing types: e.g. interactive queries on a stream, online machine learning

An Analogy



Specialized devices




New applications



Unified device

How Well Does It Work?

- 75% of users use multiple Spark components
- Competitive performance on diverse workloads





Evolution of Unification

Challenges

Spark's original “unifying abstraction” was RDDs: DAG of map/reduce steps given as Java functions on Java objects

Two challenges:

1. **Richness of optimizations**: hard to understand the computation
2. **Changing hardware limits**: I/O-bound → compute-bound

Solution: Structured APIs

New APIs that act on **structured data** (known schema) and make **semantics** of computation more visible (relational ops)

- Spark SQL engine + DataFrames + Datasets

Enable rich optimizations similar to databases and compilers

Spark SQL: Relational Data Processing in Spark

Michael Armbrust[†], Reynold S. Xin[†], Cheng Lian[†], Yin Huai[‡], Davies Liu[‡], Joseph K. Bradley[‡], Xiangrui Meng[†], Tomer Kaftan[†], Michael J. Franklin^{†‡}, Ali Ghodsi[‡], Matei Zaharia^{†*}


[†]Databricks Inc. [‡]MIT CSAIL ^{*}AMPLab, UC Berkeley

ABSTRACT

Spark SQL is a new module in Apache Spark that integrates relational processing with Spark's functional programming API. Built on our experience with Shark, Spark SQL lets Spark programmers leverage the benefits of relational processing (*e.g.*, declarative

While the popularity of relational systems shows that users often prefer writing declarative queries, the relational approach is insufficient for many big data applications. First, users want to perform ETL to and from various data sources that might be semi- or unstructured, requiring custom code. Second, users want to perform advanced analytics such as machine learning and neural networks.

Execution Steps




Example: DataFrame API

DataFrames hold rows with a known schema and offer relational operations on them through a DSL

```
> sparkR.session()  
  
> users <- sql("select * from users")  
  
> ca.users <- users[users$state == "CA"]  
            
> nrow(ca.users)      Expression AST  
  
> ca.users %>% groupBy("name") %>% avg("age")  
  
> dapply(ca.users, function(df) { df$age + rnorm(nrow(df)) }, ...)
```

What Structured APIs Enable

1. Compact binary representation
2. Optimization across operators (join order, pushdown, etc)
3. Runtime code generation



New in 2.2 & 2.3: Structured Streaming

High-level streaming API built over DataFrames

- Event time, windowing, sessions, transactional sinks
- Arbitrary or Customized Stateful Aggregation
- Stream-to-stream joins

Supports interactive & batch queries

- Ad-hoc queries via JDBC, joins with static data

Works by just incrementalizing static DataFrame queries!


Example: Batch Aggregation

```
> logs <- read.df(source = "json", path = "s3://logs")  
  
> logs %>%  
  groupBy("userid", "hour") %>%  
  avg("latency") %>%  
  write.df(source = "jdbc", path = "jdbc:mysql//...")
```

Example: Continuous Aggregation

```
> logs <- read.stream(source = "json", path = "s3://logs")  
  
> logs %>%  
  groupBy("userid", "hour") %>%  
  avg("latency") %>%  
  write.stream(source = "jdbc", path = "jdbc:mysql//...")
```

Incremental Execution




Benefits of Structured Streaming

1. Results always **consistent with a batch job** on a prefix of data
2. Same **code and libraries** run in both modes
3. Ad-hoc queries on **consistent snapshots** of state

Apache Spark: The First Unified Analytics Engine

Uniquely combines Data & AI technologies



Parquet





Databricks' Unified Analytics Platform

Unifies Data
Engineers and
Data Scientists

Unifies Data and
AI Technologies

Eliminates
infrastructure
complexity

The diagram illustrates the Databricks' Unified Analytics Platform architecture, divided into three horizontal layers:

- COLLABORATIVE WORKSPACE**: The top layer, featuring a teal gradient background. It contains icons for Data Engineers (user with gear) and Data Scientists (user with flask), along with the text "COLLABORATIVE WORKSPACE".
- DATABRICKS RUNTIME**: The middle layer, featuring an orange gradient background. It includes the text "DATABRICKS RUNTIME", the "APACHE Spark" logo, and logos for Delta, SQL, Streaming, R Studio, TensorFlow, and Databricks Learn.
- CLOUD NATIVE SERVICE**: The bottom layer, featuring a dark purple gradient background. It includes the AWS logo, the text "CLOUD NATIVE SERVICE", the Azure logo, and the Databricks logo.

Conclusion

Unified computing models are essential for usability & performance

- Same results across batch, interactive & streaming
- Powerful cross-library optimizations
- Same functionality across programming languages

In most areas of computing, we converged to unified systems

- UNIX OS, general-purpose languages, SQL, etc
- Why not in cluster computing?
- And that's **Unified Analytics Platform**
 - With **Apache Spark** at the core of **Databricks Runtime**



+





Thank You

